

<http://neyto.blogspot.co.uk/?view=magazine> – summary for docker

Outline


- **Building a container**
 - Step by step with commit
 - Entry points
- **The Dockerfile**
 - Syntax
 - Building a container
 - Supervisord
- **Mounting Volumes**
 - Adding volumes
 - Data only containers
- **Docker Hub**
 - Pushing your image

Objective

- **By the end of this session you should be able to**
 - Create a dockerfile and build our own Docker containers
 - Understand about entry points and running processes
 - Push a dockerfile to the central repository

Building Container

- **Docker containers are created from different layers**
 - Each command run creates a new layer on top of the existing structure
 - This means that common file bases can be shared between containers
- **We can build a container in one of two ways**
 - Run commands and commit them one by one
 - Write a Dockerfile - a script to build the container



4

Image from the Docker website: <https://docs.docker.com/terms/layer/>

Add a layer, commit

- **We can run commands on the base containers created by Docker**
 - If we don't have the container it will be pulled from the repository
 - Unless we are in interactive mode, each command must be able to run without any human assistance

```
# docker run centos:centos6 yum update -y
```

- This command will update centos, the container will then end
- **We can see completed containers using docker ps -l**
 - This will give us the container ID
 - To commit the changes to a container use the ID

```
# docker ps -l  
[lots of output including containerID]  
# docker commit --author="Kat" -m "msg" \ [containerID] [name:tag]
```

The naming convention is usually to use your Docker hub username and a tag to give it a version, this doesn't need to be a number. For example

```
❏ docker commit --author="Kat" -m "Updated Centos" a3e8  
kizzie/centos6:updated
```

When you use the container ID you do not need to put the full hash, usually the first four characters will suffice.

Add another layer, commit

- **This can continue until you have built everything you need**

```
# docker run kizzie/centos6:updated yum install -y java

# docker ps -l
CONTAINER ID          IMAGE                COMMAND
8bb20f239f69         kizzie/centos6:updated ...

# docker commit -author="Kat" -m "java installed"
      8bb2 kizzie/centos6:java

# docker run kizzie/centos6:java yum install -y git

... etc
```

Dockerfile

- **A Dockerfile is a script that builds the container for you**
 - It runs line by line creating new layers
 - The filename “Dockerfile” is used by convention. You can change this if you want, but you have to pass it manually to the build command

```
FROM centos:centos6
MAINTAINER Your Name

RUN yum update -y
RUN yum install -y java
Run yum install -y git
```

Syntax

- **The general form of Dockerfiles is**
 - Capitalisation of the instruction isn't required but is convention

```
INSTRUCTION arguments
```

- **Each line is executed in the order it is defined in the file**
 - You can optimise building containers by paying attention to the order of execution

From	Add
Maintainer	Copy
Run	Entrypoint
Cmd	Volume
Label	User
Expose	WorkDir
Env	OnBuild

8

The full documentation for the build files is at:

<https://docs.docker.com/reference/builder/>

Syntax Cheat Sheet

- **FROM** – The base container to use
- **MAINTAINER** – Who is responsible for the container
- **RUN** – run a command in the container
- **CMD** – One in a file, argument to pass the default shell on startup
- **LABEL** – Add meta data to the image, key value pairs
- **EXPOSE** – Ports to allow access through
- **ENV** – Environment variables
- **ADD** – Copy files to the container, can be a URL, will decompress tar
- **COPY** – Same as add but without the decompress or URL
- **ENTRYPOINT** – Command to be run when the container is started
- **VOLUME** – Mount a directory as a volume for the container
- **USER** – Sets the current user from this point forward
- **WORKDIR** – changes the working directory from this point forward
- **ONBUILD** – Trigger for a build event, not passed forward to children

Building the container

- **Once you are happy with your file we can trigger a build with**

- The dot at the end is important!
 - This is the context the docker file builds in
 - Important for copy / add
- The build will look for something called "Dockerfile" in the current directory
- If you want to change the file name use the --file flag

```
# docker build -t [imagename] .  
# docker build -t [imagename] --file="fileName" .
```

- **Each intermediate step of the build is given an ID**

- If something fails you can jump into the container at a given point with the ID and take a look around with bash

Supervisord

- **We can't just start a container and let it run**
 - Remember – it's not the same as a virtual machine!
- **We either need to specify the command to run each time, or use an entry point**
 - Both CMD and ENTRYPOINT do similar jobs
 - Best practice is to use entrypoint and use cmd for arguments
 - However, if the script or commands ends, the container stops
- **Supervisord as a solution:**
 - Use a supervisord as the entry point
 - The supervisor will run any commands in it's config file -
/etc/supervisor/conf.d/supervisord.conf

```
[supervisord]
nodaemon=true

[program:mysql]
command=service mysql start
```

Mounting Volumes

- **Containers should be able to be created and destroyed at will**
 - The data inside a container is lost when the container is removed
 - Every time a file is changed it generates another layer in the container
 - For persistent data we should be using volumes
 - Volumes persist as long as there is a container referencing it
- **A volume is a directory mounted from outside the container**
 - It can be on the local system
 - Or from another container
- **File systems can be shared between containers**
 - Files contained in the volume are linked to the container

12

See <http://container-solutions.com/understanding-volumes-docker/> for a very well put together explanation of how the layers file system and volumes interact.

Using the local file system

- **The -v flag will mount a volume**
 - Creates a directory /test inside the volume
 - Stored on the host system in /var/lib/docker/volumes/
 - Hard to find our files again...

```
# docker run -it -v /test ubuntu bash
```

- **Better to give it a location to use for the host directory**

```
# docker run -v hostDirectory:containerDirectory ...
```

```
# docker run -it -v /testHost:/testVolume/ ubuntu bash
root@53e220260718:/testVolume# echo
'hello world' > /testVolume/hello
root@53e220260718:/testVolume# exit
Exit
root@user-VirtualBox# cat /testHost/hello
hello world
```

Data only containers

- **Create a container that doesn't run, but acts as a host for the volumes**
 - Access it via the `--volumes-from` flag
 - Docker now controls the data volumes in an accessible way as well as the running containers

```
# docker run
  -d
  -v /var/lib/mysql
  --name data-container-mysql
  mysql
  echo data-only container for mysql
```

- Creates a container using the `mysql` image, runs the `echo` command and then stops

Volumes From

- **To connect other containers we use the `--volumes-from` flag**

```
$ docker run
  -d
  --volumes-from data-container-mysql
  -e MYSQL_ROOT_PASS="password"
  mysql
```

- This will mount the directories from the data-container-mysql image to this one
- We can create a database inside this container, stop the container, create a new one using the `--volumes-from` command and the database will still be there
- **The container can then be stored in the registry with the current data**

Docker hub - <https://hub.docker.com/>

- **The Docker hub acts as a repository for all the containers built by users**
 - There are a set of official containers – they do not have a username associated
 - Best practice is to always start by building from these containers
- **You can search the repository and pull down containers**

```
# docker search [term]
# docker pull [imagename]
```

- **Docker hub can even be linked directly to git and other services**
 - Push new code to the repository
 - The hub clones the repo and builds the container based on a Dockerfile present in the root directory

Pushing our container

- **To push one of our containers to the hub we need to ensure we have named it correctly**
 - Sign up for an account
 - Your build name should include your username

```
username/imagename:tag
```

- **Then we can push this directly to docker**
 - The container will be uploaded layer by layer
 - It will check for existing layers already saved

```
[root@localhost centos_java_helloworld]# docker push kizzie/centos:java
The push refers to a repository [kizzie/centos] (len: 1)
Sending image list
Please login prior to push:
Username: kizzie
Password:
Email: katrina.mcivor@qa.com
Login Succeeded
The push refers to a repository [kizzie/centos] (len: 1)
Sending image list
Pushing repository kizzie/centos (1 tags)
f1b10cd84249: Image already pushed, skipping
b9aeaeab5e17: Image already pushed, skipping
d40a940f40c6: Pushing [====>
```

```
] 24.41 MB/204.3 MB 9m7s
```

Exercise

- **Write and use some Dockerfiles**
 - Build an existing Dockerfile for a project
 - Write a Dockerfile to build and run the hello-scalatra project
 - Create a tomcat container for the hello-scalatra project

Summary

- **Building a container**
 - Step by step with commit
 - Entry points
- **The Dockerfile**
 - Syntax
 - Building a container
 - Supervisord
- **Mounting Volumes**
 - Adding volumes
 - Data only containers
- **Docker Hub**
 - Pushing your image

Best Practices

- **Containers should be ephemeral**
 - If something has stopped working, kill it and start a new one
 - Don't store information in the container – mount volumes instead
- **Use a .dockerignore file**
 - Don't copy over directories and files that are not important to the container
 - Containers should be as small as possible
- **Avoid installing unnecessary packages**
 - Again, we're trying to be small and fast!

https://docs.docker.com/articles/dockerfile_best-practices/

Best Practices

- **Run one process per container**
 - We can link containers together if we need apache AND mysql
 - This way we can replace the mysql container for another seamlessly
 - Separation of concerns
- **Minimise layers**
 - Readability vs Maintainability here
 - More layers means more to pull down from the repository each time
- **Sort multi-line arguments**
 - Formatting the file to be readable
 - Use the backslash to separate commands to new lines

```
RUN apt-get update && apt-get install -y \  
    bzip \br/>    cvs \br/>    git
```

21

Best Practices

- **Build from the cache**
 - We can turn off the cache with the `--no-cache` flag
 - But why run `yum update` twice?
 - If files are added or copied a comparison is made to see if they have changed, if so it will spawn a new container
 - It will not check git repositories if you clone them directly!
- **FROM – Use official repositories**
 - Your container will always run and be up to date
 - No one can sneakily add a backdoor
- **RUN – Format long commands to make them readable**
 - Use the backslashes to separate over lines
 - Don't run `upgrade` as this will likely fail due to privileges

Best Practices

- **CMD – only used to run software in the image along with arguments**
 - Should use the form ["exec", "arg1", "arg2"]
 - CMD ["apache", "-DFOREGROUND"] is recommended for a service based image
- **EXPOSE – only expose what you need to**
 - Use the traditional ports for applications
 - If we have three containers running something on 8080 we can just redirect their output on the host
- **ENV – update path variables**
 - Means that your cmd operations can be simpler

Best Practices

- **ADD and COPY**

- They do similar things
- If you don't need to get something from a URL or to decompress it then use COPY
- Copy files individually so the cache will only invalidate a container when one file changes
- You are discouraged from using ADD to get files from URLs. Use wget or curl instead

- **ENTRYPOINT**

- Best practice is to use this for the command to be run when the container starts
- Use CMD as default flags

Best Practices

- **VOLUME**
 - Expose database storage areas
 - Store files generated
 - Config storage
- **USER**
 - You're already super user in the container
 - You can add users and change to non-root users with this
 - Avoid installing sudo – unpredictable results
 - Avoid switching user to often – minimise layers
- **WORKDIR**
 - Use absolute paths
 - Use to avoid "RUN cd something & do something else & cd .."